

Bucketing in Spark

Bucketing

Bucketing is a data organization technique in Apache Spark that divides datasets into manageable chunks based on hash values. This technique significantly improves performance for operations like:

1. Filtering
2. Aggregations (group by)
3. Joins

```
df.filter(product_id = x)
```


1. as-it-is: No organization
Pro: simple
Con: full table scan
use: small datasets only
2. Partitioning
Pro: good for low cardinality column
Con: create small file problem
use: date-based cols, categories
3. Bucketing
Pro: Avoids shuffle: optimal for high cardinality columns
Con: requires upfront planning
use: high cardinality join/filter cols

Orders Table:

order_id	product_id	customer_id	quantity	total_amt	order_date
1	22	105	10	104	2023-03-10
2	19	542	12	114	2023-03-10
3	11	199	4	2567	2023-03-10
4	87	76	10	1268	2023-03-10
5	9	225	6	1136	2023-03-10
6	24	980	11	658	2023-03-10
7	22	14	55	572	2023-03-10
8	10	5	12	7768	2023-03-10

```
bucket_assignment =
hash(bucketing_column) % number_of_buckets
h(product_id) % 4
```

4 Buckets



1. filter

Without Bucketing:

Query: WHERE product_id = 22

Problem: Scans ALL records

With Bucketing

Query: WHERE product_id = 22

Optimization:

1. Calculate: $22 \% 4 = 2$
2. Read ONLY Bucket 2
3. Skip Buckets 0, 1, 3

Result: Reduced search space by 75%!**This is called "Bucket Pruning"**

2. Joins

1. as-it-is

→ make sure that same key is in the same partition

3 step operation: shuffle | sort | merge

Step 1: SHUFFLE - Redistribute data across nodes (COSTLY!)

Step 2: SORT - Sort by join key

Step 3: MERGE - Perform the join



Physical Plan


```
Exchange/Shuffle (Orders) ←— EXPENSIVE
  ↓
Sort (Orders)
  ↓
Exchange/Shuffle (Products) ←— EXPENSIVE
  ↓
Sort (Products)
  ↓
SortMergeJoin
```

2. Partitioning

- small file problem
- product_id is a high cardinality column

3. Bucketing






Executor
1

Executor
2

Executor
3

Executor
4



Join Scenarios with Bucketing

Scenario Matrix

Dataset 1	Dataset 2	Shuffle Required?	Performance
Bucketed (X buckets, col A)	Bucketed (X buckets, col A)	✗ No	★★★ Excellent
Bucketed (X buckets, col A)	Bucketed (Y buckets, col A)	⚠ Partial (one dataset)	★★ Good
Bucketed (X buckets, col A)	Bucketed (X buckets, col A), but JOIN on col B	✓ Full Shuffle	★ Poor
Not Bucketed	Not Bucketed	✓ Full Shuffle	★ Poor

3. GroupBy

```
SQL
SELECT product_id, SUM(total_amt) as total_sales
FROM orders
GROUP BY product_id
```

Without Bucketing:

Step 1: Local/Partial Aggregation
 Step 2: SHUFFLE (Exchange HashPartitioning) ← EXPENSIVE
 Step 3: Global Aggregation

With Bucketing:

Step 1: Local/Partial Aggregation (per bucket on same executor)
 Step 2: Global Aggregation (no shuffle needed!)
 SHUFFLE ELIMINATED!

Why? All rows with the same `product_id` are already in the same bucket on the same executor.

Determining Optimal Bucket Count

Formula

Number of Buckets = Dataset Size (MB) / Optimal Bucket Size (MB)

where Optimal Bucket Size = 128-200 MB

Example Calculation

Dataset Size = 1000 MB (1 GB)

Optimal Bucket Size = 200 MB

Number of Buckets = 1000 / 200 = 5 buckets

Estimating Dataset Size

When data is not yet written to disk, use this formula:

Dataset Size (MB) = $(N \times V \times W) / (1024^2)$

where:

N = Number of records

V = Number of variables (columns)

W = Average width in bytes per variable

Suggested Readings:

1. https://umbertogriffo.gitbook.io/apache-spark-best-practices-and-tuning/parallelism/sparksqlexportpartitions_draft
2. <https://medium.com/globant/how-to-solve-a-large-number-of-small-files-problem-in-spark-21f819eb36d3>