

# 01 Databricks Lakeflow Declarative

## Comprehensive Step-by-Step Guide: Databricks Lakeflow Declarative Pipeline with Medallion Architecture

### Project Overview

This guide walks you through building an end-to-end data engineering solution using Databricks Lakeflow Declarative Pipelines, implementing a Medallion Architecture (Bronze, Silver, Gold) with automated incremental data loading.

---

## 1. Prerequisites and Data Preparation

### Data Files Required

You need the following CSV files:

#### Fact Table:

- **fact\_sales.csv** - Contains transaction data with columns:
  - sale\_id, order\_date, customer\_id, product\_id, quantity, discount, region\_id, channel, promo\_code
  - Approximately 50,000 records
  - Note: Contains null values in quantity, channel, and promo\_code columns

#### Dimension Tables:

- **customer.csv** - Customer details:
  - customer\_id, first\_name, last\_name, email, join\_date, VIP
  - Contains some missing email values
- **product.csv** - Product information:
  - product\_id, product\_name, category, price, in\_stock
- **region.csv** - Region details:
  - region\_id, region\_name, country

#### Additional File for Testing:

- **fact\_sales\_2025.csv** - For testing incremental load (approximately 2,526 records)

### Account Setup

- Sign up for Databricks (free version available)
- Ensure you have access to Azure Databricks workspace

## 2. Unity Catalog and Schema Setup

### Step 2.1: Create Unity Catalog

1. Navigate to **Catalog** in Databricks workspace
2. Click **Create Catalog**
3. Name: `lakeflow_dlt`
4. Click **Create**

### Step 2.2: Create Four Schemas

Create the following schemas within the `lakeflow_dlt` catalog:

#### a) Landing Zone Schema:

1. Click on catalog name `lakeflow_dlt`
2. Click **Create Schema** (top right)
3. Name: `landing_zone`
4. Click **Create**

#### b) Bronze Schema:

1. Click **Create Schema**
2. Name: `bronze`
3. Click **Create**

#### c) Silver Schema:

1. Click **Create Schema**
2. Name: `silver`
3. Click **Create**

#### d) Gold Schema:

1. Click **Create Schema**
2. Name: `gold`
3. Click **Create**

**Result:** You should now have:

- `lakeflow_dlt.landing_zone`
- `lakeflow_dlt.bronze`
- `lakeflow_dlt.silver`
- `lakeflow_dlt.gold`

### 3. Volume Creation and Data Ingestion

#### Step 3.1: Create Managed Volume

1. Navigate to `lakeflow_dlt` catalog
2. Select `landing_zone` schema
3. Click **Create Volume**
4. Name: `fact_and_dimension_files`
5. Type: **Managed Volume**
6. Click **Create**

#### Step 3.2: Create Directory Structure

Create separate directories for each file to avoid conflicts:

##### a) Fact Sales Directory:

1. Inside the volume, click **Create Directory**
2. Name: `fact_sales`
3. Click **Create**
4. Double-click to open the directory
5. Click **Upload**
6. Select `fact_sales.csv`
7. Upload the file

##### b) Customers Directory:

1. Navigate back to main volume directory
2. Click **Create Directory**
3. Name: `dim_customers`
4. Click **Create**
5. Open directory and upload `customer.csv`

##### c) Products Directory:

1. Navigate back to main volume directory
2. Click **Create Directory**
3. Name: `dim_products`
4. Click **Create**
5. Open directory and upload `product.csv`

##### d) Regions Directory:

1. Navigate back to main volume directory
2. Click **Create Directory**
3. Name: `dim_regions`
4. Click **Create**

5. Open directory and upload `region.csv`

**Important:** Each file must be in its own directory for proper AutoLoader functionality.

---

## 4. Lakeflow Pipeline Creation

### Step 4.1: Initialize Pipeline

1. Navigate to **Jobs and Pipelines** (or **Workflows**)
2. Click **Build ETL Pipeline**
3. Pipeline Name: `databricks_declarative_lakeflow_pipeline`
4. Press **Enter** to save

### Step 4.2: Configure Pipeline Settings

1. **Unity Catalog:** Select `lakeflow_dlt`
2. **Schema:** Select `landing_zone`
3. **Start Option:** Choose **Empty File** (not sample code)
4. **Language:** Select **SQL**
5. Click **Select**

### Step 4.3: Understand Pipeline Interface

- **Left Panel:** Pipeline assets and transformations
- **Center:** Code editor
- **Right Panel:** Pipeline graph visualization
- **Bottom:** Tables, performance metrics, and results
- **Top Right:** Settings, dry run, and run pipeline buttons

### Step 4.4: Enable New Pipeline Editor

1. Toggle **Lakeflow pipelines editor** ON (recommended)
  2. This provides better visualization and debugging capabilities
- 

## 5. Bronze Layer - Data Ingestion

### Step 5.1: Create Bronze Transformation File

1. In the transformations folder, you'll see `my_transformation.sql`
2. Right-click or click kebab menu
3. Select **Rename**
4. New name: `bronze_transformations.sql`
5. Click **OK**

### Step 5.2: Write Bronze Layer Code

#### Fact Sales Ingestion:

```
-- Ingest data from volume to bronze schema

CREATE STREAMING LIVE TABLE bronze.fact_sales
COMMENT "Raw fact sales from volume to bronze schema"
AS SELECT * FROM CLOUD_FILES(
  '/Volumes/lakeflow_dlt/landing_zone/fact_and_dimension_files/fact_sales',
  'csv',
  map('header', 'true')
);
```

SQL

**Products Ingestion:**

```
CREATE STREAMING LIVE TABLE bronze.products
COMMENT "Raw products from volume to bronze schema"
AS SELECT * FROM CLOUD_FILES(
  '/Volumes/lakeflow_dlt/landing_zone/fact_and_dimension_files/dim_products',
  'csv',
  map('header', 'true')
);
```

SQL

**Customers Ingestion:**

```
CREATE STREAMING LIVE TABLE bronze.customers
COMMENT "Raw customers from volume to bronze schema"
AS SELECT * FROM CLOUD_FILES(
  '/Volumes/lakeflow_dlt/landing_zone/fact_and_dimension_files/dim_customers',
  'csv',
  map('header', 'true')
);
```

SQL

**Regions Ingestion:**

```
CREATE STREAMING LIVE TABLE bronze.regions
COMMENT "Raw regions from volume to bronze schema"
AS SELECT * FROM CLOUD_FILES(
  '/Volumes/lakeflow_dlt/landing_zone/fact_and_dimension_files/dim_regions',
  'csv',
  map('header', 'true')
);
```

SQL

**Step 5.3: Get Volume Paths**

Instead of typing paths manually:

1. Open **Catalog** panel on the right
2. Navigate to: `lakeflow_dlt` > `landing_zone` > `fact_and_dimension_files`
3. Expand to see directories
4. Click the **double chevron (>>)** icon next to directory name
5. Path auto-populates in your query
6. Remove trailing forward slash if present

### Step 5.4: Validate with Dry Run

1. Click **Dry Run** button (top right)
2. This validates schema and syntax without executing
3. Check for errors in the output panel
4. Fix any issues (common: duplicate names, typos)
5. Wait for "Run completed" success message

### Step 5.5: Execute Pipeline

1. Click **Run** button
2. Wait for pipeline execution
3. Monitor progress in pipeline graph
4. Green checkmarks indicate success
5. View tables in **All Tables** section

### Step 5.6: Verify Results

1. Click **Maximize** to view full pipeline graph
2. Verify all four streaming tables created:
  - `bronze.fact_sales` (50,000 records)
  - `bronze.customers` (5,000 records)
  - `bronze.products` (500 records)
  - `bronze.regions` (10 records)
3. Click on any table to view:
  - **Preview:** Sample data
  - **Table Metrics:** Row counts, duration
  - **Columns:** Data types (all string initially)

## 6. Silver Layer - Data Cleaning

### Step 6.1: Create Silver Transformation File

1. Click **Create File** in transformations folder
2. Language: **SQL**
3. Name: `silver_transformations.sql`
4. Dataset Type: **Streaming Table**
5. Click **Create**

## Step 6.2: Write Data Cleaning Code

### Fact Sales Cleaning:

```
-- Fact sales data cleaning

CREATE STREAMING LIVE TABLE silver.fact_sales
COMMENT "Fact sales data cleaning"
AS SELECT
    CAST(sale_id AS INT) AS sale_id,
    TO_DATE(order_date, 'dd/MM/yyyy') AS order_date,
    CAST(customer_id AS INT) AS customer_id,
    CAST(product_id AS INT) AS product_id,
    CAST(quantity AS INT) AS quantity,
    CAST(discount AS INT) AS discount,
    CAST(region_id AS INT) AS region_id,
    CAST(channel AS STRING) AS channel,
    CAST(promo_code AS STRING) AS promo_code
FROM STREAM(lakeflow_dlt.bronze.fact_sales);
```

SQL

### Products Cleaning:

```
CREATE STREAMING LIVE TABLE silver.products
COMMENT "Products data cleaning"
AS SELECT
    CAST(product_id AS INT) AS product_id,
    CAST(product_name AS STRING) AS product_name,
    CAST(category AS STRING) AS category,
    CAST(price AS INT) AS price,
    CAST(in_stock AS INT) AS in_stock
FROM STREAM(lakeflow_dlt.bronze.products);
```

SQL

### Customers Cleaning:

```
CREATE STREAMING LIVE TABLE silver.customers
COMMENT "Customers data cleaning"
AS SELECT
    CAST(customer_id AS INT) AS customer_id,
    CAST(first_name AS STRING) AS first_name,
    CAST(last_name AS STRING) AS last_name,
    CAST(email AS STRING) AS email,
    TO_DATE(join_date, 'dd/MM/yyyy') AS join_date,
    CAST(VIP AS STRING) AS VIP
FROM STREAM(lakeflow_dlt.bronze.customers);
```

SQL

## Regions Cleaning:

```
CREATE STREAMING LIVE TABLE silver.regions
COMMENT "Regions data cleaning"
AS SELECT
    CAST(region_id AS INT) AS region_id,
    CAST(region_name AS STRING) AS region_name,
    CAST(country AS STRING) AS country
FROM STREAM(lakeflow_dlt.bronze.regions);
```

SQL

## Step 6.3: Important - Add STREAM() Wrapper

**Critical Fix:** Wrap the source path with **STREAM()** to enable incremental processing:

```
FROM STREAM(lakeflow_dlt.bronze.fact_sales)
```

SQL

This prevents the error: "Cannot create a streaming table append flow from a batch query"

## Step 6.4: Run and Verify

1. Click **Dry Run** to validate
2. Fix any errors (common: duplicate table names, missing STREAM wrapper)
3. Click **Run**
4. Verify 8 total streaming tables (4 bronze + 4 silver)
5. Check data types in silver tables are properly converted

## 7. Silver Layer - Materialized Views with Constraints

### Step 7.1: Create Cleaned Sales Data File

1. Click **Create File**
2. Language: **SQL**
3. Name: **cleaned\_sales\_data.sql**
4. Dataset Type: **Materialized View**
5. Click **Create**

## Step 7.2: Write Materialized View with Constraints

SQL

```

CREATE OR REFRESH MATERIALIZED VIEW silver.cleaned_sales_data
(
    -- Define data quality constraints
    CONSTRAINT sales_quantity_check EXPECT (quantity IS NOT NULL) ON VIOLATION DROP ROW,
    CONSTRAINT sales_channel_check EXPECT (channel IS NOT NULL) ON VIOLATION DROP ROW,
    CONSTRAINT sales_promo_code_check EXPECT (promo_code IS NOT NULL) ON VIOLATION DROP ROW,
    CONSTRAINT customer_email_check EXPECT (email IS NOT NULL) ON VIOLATION DROP ROW
)
COMMENT "Cleaned sales data"
AS SELECT
    -- Fact sales columns
    FS.sale_id,
    FS.order_date,
    FS.customer_id,
    FS.product_id,
    FS.quantity,
    FS.discount,
    FS.region_id,
    FS.channel,
    FS.promo_code,
    -- Customer columns (exclude primary key)
    C.first_name,
    C.last_name,
    C.email,
    C.join_date,
    C.VIP,
    -- Product columns (exclude primary key)
    P.product_name,
    P.category,
    P.price,
    P.in_stock,
    -- Region columns (exclude primary key)
    R.region_name,
    R.country,
    -- Calculated column
    FS.quantity * P.price AS revenue
)

```

```

FROM lakeflow_dlt.silver.fact_sales FS

LEFT JOIN lakeflow_dlt.silver.customers C
  ON FS.customer_id = C.customer_id

LEFT JOIN lakeflow_dlt.silver.products P
  ON FS.product_id = P.product_id

LEFT JOIN lakeflow_dlt.silver.regions R
  ON FS.region_id = R.region_id;

```

## Step 7.3: Understanding Constraints

### Constraint Syntax:

`CONSTRAINT <constraint_name> EXPECT (<condition>) ON VIOLATION <action>`

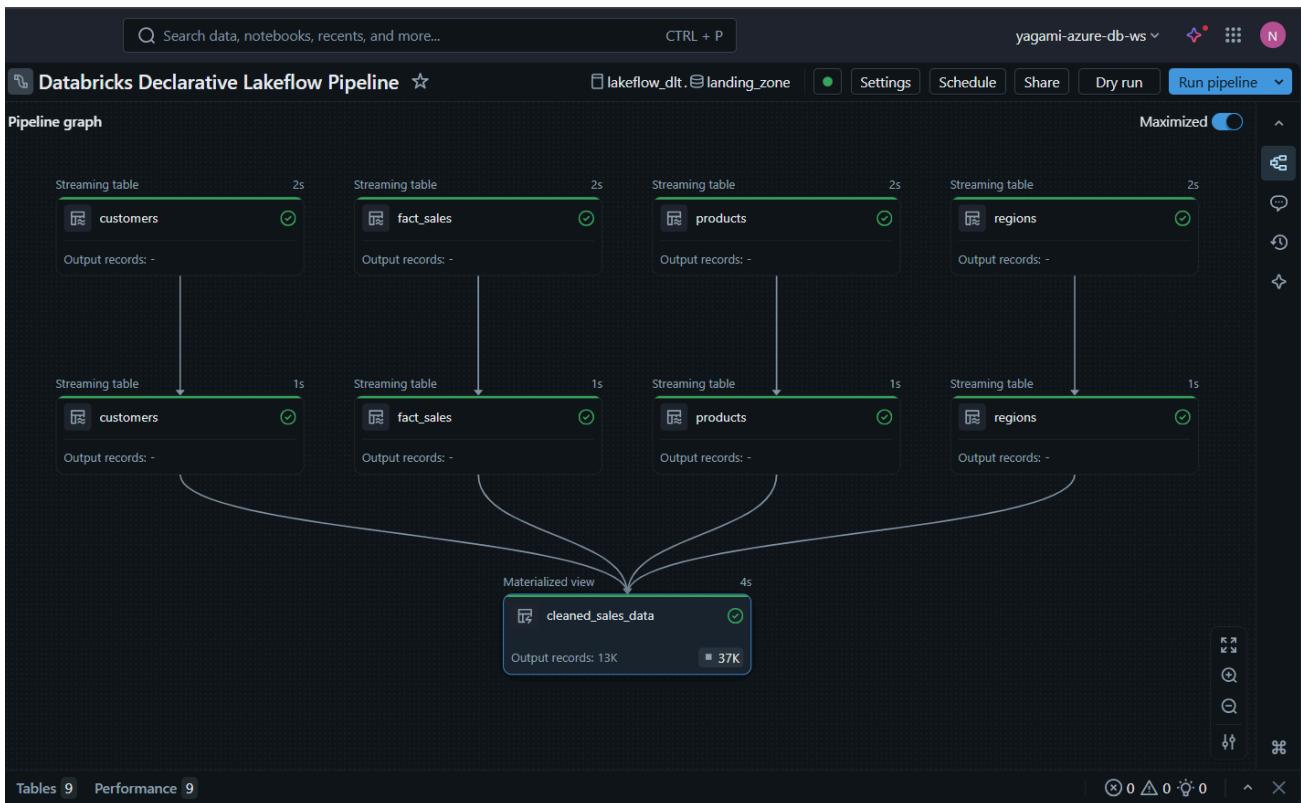
SQL

### Actions:

- **DROP ROW** - Removes rows that violate constraint
- **FAIL** - Stops pipeline on violation
- **QUARANTINE** - Moves bad data to separate table

## Step 7.4: Run and Validate

1. Click **Dry Run**
2. Click **Run**
3. View pipeline graph - materialized view shown with turbo icon
4. Check **Table Metrics** to see:
  - Total records written: ~12,935 (25.9%)
  - Total records dropped: ~37,065 (74.1%)
  - Breakdown by constraint



## Step 7.5: Analyze Data Quality Results

Click on table metrics to see:

- **sales\_promo\_code\_check:** 36,533 rows dropped
- **sales\_channel\_check:** 999 rows dropped (2%)
- **customer\_email\_check:** 971 rows dropped
- **sales\_quantity\_check:** 44 rows dropped (<1%)

**Note:** In production, 74% drop rate is high. This demonstrates the feature; adjust constraints to business needs.

## 8. Gold Layer - Business Aggregations

### Step 8.1: Create Gold Transformations File

1. Click **Create File**
2. Language: **SQL**
3. Name: **gold\_transformations.sql**
4. Dataset Type: **Materialized View**
5. Click **Create**

### Step 8.2: Write Business Aggregate Queries

#### Aggregate 1: Category Sales Summary

SQL

```
CREATE OR REFRESH MATERIALIZED VIEW gold.category_sales_summary
AS SELECT
    category,
    YEAR(order_date) AS year,
    SUM(revenue) AS total_revenue
FROM lakeflow_dlt.silver.cleaned_sales_data
GROUP BY category, YEAR(order_date)
ORDER BY category, year;
```

## Aggregate 2: Revenue by Customer in Each Region (Ranked)

SQL

```
CREATE OR REFRESH MATERIALIZED VIEW
gold.revenue_by_customers_in_each_region_by_ranking
AS SELECT
    region_name,
    customer_id,
    first_name,
    last_name,
    SUM(revenue) AS total_revenue,
    RANK() OVER (PARTITION BY region_name ORDER BY SUM(revenue) DESC) AS revenue_rank
FROM lakeflow_dlt.silver.cleaned_sales_data
GROUP BY region_name, customer_id, first_name, last_name
ORDER BY region_name, revenue_rank;
```

## Aggregate 3: Customer Lifetime Value Estimation

SQL

```
CREATE OR REFRESH MATERIALIZED VIEW gold.customer_lifetime_value_estimation
AS SELECT
    customer_id,
    first_name,
    last_name,
    email,
    VIP,
    COUNT(DISTINCT sale_id) AS total_orders,
    SUM(quantity) AS total_items_purchased,
    SUM(revenue) AS lifetime_value,
    AVG(revenue) AS average_order_value,
    MIN(order_date) AS first_purchase_date,
    MAX(order_date) AS last_purchase_date,
    DATEDIFF(MAX(order_date), MIN(order_date)) AS customer_tenure_days
FROM lakeflow_dlt.silver.cleaned_sales_data
GROUP BY customer_id, first_name, last_name, email, VIP
ORDER BY lifetime_value DESC;
```

## Step 8.3: Execute and Verify

1. Click **Dry Run**
2. Click **Run**
3. View complete pipeline graph showing full medallion architecture
4. Verify three gold layer materialized views:
  - `gold.category_sales_summary` (~56 records)
  - `gold.revenue_by_customers_in_each_region_by_ranking` (~52 records)
  - `gold.customer_lifetime_value_estimation` (~4,600 records)

## Step 8.4: Customize Pipeline Graph View

1. Click **Maximize** for full view
  2. Use **Fit to View** to adjust zoom
  3. Change **Graph Orientation**:
    - Horizontal Layout (default)
    - **Vertical Layout** (recommended for medallion flow)
  4. Toggle **Mini Map** on/off as needed
  5. Click **Refresh** in catalog to see all tables
- 

## 9. Scheduling and Automation

### Step 9.1: Configure Pipeline Schedule

1. Click **Schedule** button in pipeline interface
2. **Schedule Type:** Choose frequency
  - **Option A:** Advanced - Every 3 minutes (for demo/testing)
  - **Option B:** Daily, Weekly, or Custom cron
3. **Schedule Name:** Keep default or customize
4. **Time Zone:** Verify correct timezone
5. Click **Create Schedule**

### Step 9.2: Verify Schedule

1. Schedule status shows **Active**
2. Pipeline will run automatically at specified intervals
3. Monitor execution in **Jobs and Pipelines**

### Step 9.3: Compute Configuration (Optional)

By default, pipeline uses:

- **Serverless Compute** (recommended)
- Auto-scaling based on workload
- No manual cluster management required

To modify:

1. Click **Settings** (gear icon)
  2. Adjust compute settings if needed
  3. Review JSON/YAML configuration
- 

## 10. Dashboard Creation

### Step 10.1: Create New Dashboard

1. Navigate to **Dashboards**
2. Click **Create Dashboard**
3. Name: **Category Sales Summary** (or your preference)

### Step 10.2: Add Data Source

1. Click **Add > Visualization**
2. Select **Query Data**
3. Navigate to: **lakeflow\_dlt > gold > category\_sales\_summary**
4. Click table to select
5. Wait for data to load (serverless warehouse starts automatically)

### Step 10.3: Create Visualizations

#### Visualization 1: Total Revenue Counter

1. Click **Add Visualization**
2. In AI assistant prompt: "Show total revenue as counter"
3. Press **Enter**
4. AI generates counter showing total revenue (e.g., £2.06M)

#### Visualization 2: Revenue by Category Bar Chart

1. Click **Add Visualization**
2. Prompt: "Revenue by category using bar chart"
3. Press **Enter**
4. AI generates horizontal or vertical bar chart

#### Visualization 3: Year Filter

1. Click **Add Filter**
2. Select **Add Field**
3. Choose **year** column
4. Shows available years (2018-2024)
5. Users can filter dashboard by year

### Step 10.4: Arrange Dashboard

1. Drag and resize visualizations
2. Position counter at top

3. Place bar chart prominently
4. Add year filter for interactivity

## Step 10.5: Publish Dashboard

1. Click **Publish** button
2. Dashboard is now live and accessible
3. Share URL with stakeholders



## 11. Testing Incremental Load

### Step 11.1: Prepare Test Data

You should have `fact_sales_2025.csv` ready with:

- Approximately 2,526 records
- Same schema as original fact\_sales.csv
- Data for year 2025

### Step 11.2: Upload New Data

1. Navigate to **Catalog**
2. Go to: `lakeflow_dlt > landing_zone > fact_and_dimension_files`
3. Open `fact_sales` directory
4. Click **Upload to this volume**
5. Click **Browse**
6. Select `fact_sales_2025.csv`
7. Click **Upload**

### Step 11.3: Monitor Scheduled Execution

1. Navigate to **Jobs and Pipelines**
2. Find your scheduled pipeline
3. Wait for next scheduled run (3 minutes if using that interval)

4. Monitor execution status
5. Pipeline automatically detects new file via AutoLoader

### Step 11.4: Verify Incremental Load

#### Check Pipeline Graph:

1. Open pipeline after execution
2. Record counts should increase
3. `bronze.fact_sales`: ~52,500 records (50,000 + 2,526 - duplicates removed)
4. Downstream tables update accordingly

#### Check Dashboard:

1. Refresh your dashboard
2. Year filter now shows **2025** option
3. Select 2025 to see new data
4. Total revenue includes 2025 transactions (£935 for 2025)
5. Bar chart shows 2025 category breakdown

#### Verify in Catalog:

1. Go to `gold.category_sales_summary`
2. Click **Sample Data**
3. Scroll to find 2025 records
4. Confirm new year data is present

### Step 11.5: Validate AutoLoader Functionality

AutoLoader automatically:

- Detects new files in volume directories
- Processes only new/changed files (incremental)
- Maintains state to avoid reprocessing
- Handles schema evolution (if configured)
- Provides exactly-once processing guarantees

## Additional Features and Best Practices

### Pipeline Management

#### Version History:

- Click **Version History** to see all changes
- Track who made changes and when
- Revert to previous versions if needed

## Collaboration:

- Click **Comments** icon to add annotations
- Team members can review and discuss
- Tag colleagues for feedback

## Error Handling:

- View **Flows** tab to see execution details
- Check **Expectations** section for constraint violations
- Review **Failed** records for debugging

## Performance Optimization

### 1. Partition Strategy:

```
CREATE STREAMING LIVE TABLE bronze.fact_sales
PARTITIONED BY (YEAR(order_date))
AS SELECT * FROM ...
```

SQL

### 2. Z-Ordering:

```
CREATE OR REFRESH MATERIALIZED VIEW gold.customer_lifetime_value_estimation
CLUSTER BY (lifetime_value)
AS SELECT ...
```

SQL

### 3. Auto Optimize:

Enable in pipeline settings for automatic file compaction

## Data Quality Enhancements

### Custom Expectations:

```
CONSTRAINT price_validation
EXPECT (price > 0)
ON VIOLATION DROP ROW

CONSTRAINT date_range_check
EXPECT (order_date BETWEEN '2018-01-01' AND CURRENT_DATE())
ON VIOLATION QUARANTINE
```

SQL

**Quarantine Tables:** Use **ON VIOLATION QUARANTINE** to review bad data later instead of dropping

## Security Configuration

### 1. Table ACLs:

```
GRANT SELECT ON TABLE gold.category_sales_summary TO `analysts@company.com`;
```

SQL

## 2. Row-Level Security:

```
CREATE OR REFRESH MATERIALIZED VIEW gold.regional_sales
AS SELECT * FROM cleaned_sales_data
WHERE region_name = current_user_region();
```

SQL

# Troubleshooting Common Issues

## Issue 1: "Cannot redefine dataset"

**Cause:** Duplicate table names in same schema

**Solution:** Ensure unique names for each streaming table/materialized view

## Issue 2: "Cannot create streaming table from batch query"

**Cause:** Missing `STREAM()` wrapper

**Solution:** Wrap source with `STREAM(source_table)`

## Issue 3: AutoLoader not picking up new files

**Cause:** File format or path mismatch

**Solution:**

- Verify file format matches (csv, parquet, etc.)
- Check directory path is correct
- Ensure file structure matches original

## Issue 4: High constraint violation rate

**Cause:** Overly strict constraints or data quality issues

**Solution:**

- Review constraint logic
- Analyze source data quality
- Consider `QUARANTINE` instead of `DROP ROW`

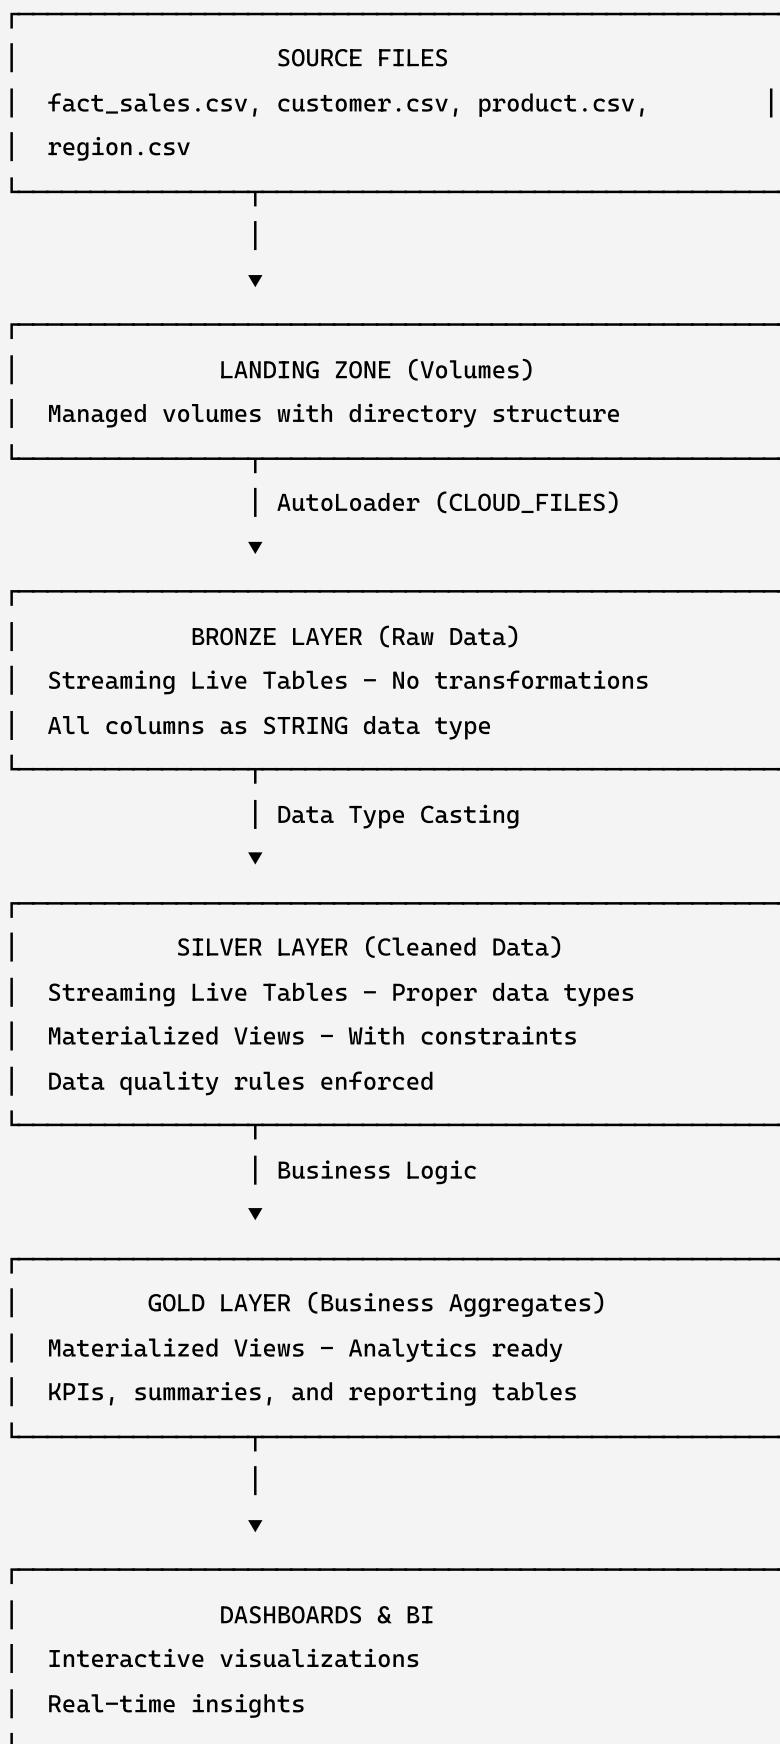
## Issue 5: Dashboard not updating

**Cause:** Cached results or warehouse stopped

**Solution:**

- Refresh dashboard manually
- Check if serverless warehouse is running
- Verify pipeline executed successfully

## Project Architecture Summary



## Key Learnings and Takeaways

1. **Declarative Pipelines:** Define what you want, not how to do it
  2. **AutoLoader:** Automatically handles incremental data ingestion
  3. **Medallion Architecture:** Clear separation of raw, cleaned, and aggregated data
  4. **Data Quality:** Built-in constraints ensure data integrity
  5. **Streaming Tables:** Enable real-time data processing
  6. **Materialized Views:** Pre-computed results for faster queries
  7. **Serverless Compute:** No cluster management required
  8. **End-to-End Automation:** From ingestion to visualization
- 

## Next Steps and Enhancements

1. **Add more data sources:** Kafka, Kinesis, S3 events
  2. **Implement Change Data Capture (CDC):** Track data changes
  3. **Add data lineage tracking:** Understand data flow
  4. **Integrate with ML workflows:** Use gold tables for training
  5. **Set up alerting:** Notify on pipeline failures or data quality issues
  6. **Implement slowly changing dimensions (SCD):** Track historical changes
  7. **Add data retention policies:** Manage storage costs
  8. **Create additional dashboards:** For different business units
- 

## Resources

- **Databricks Documentation:** <https://docs.databricks.com>
  - **Unity Catalog Guide:** <https://docs.databricks.com/unity-catalog>
  - **Delta Live Tables:** <https://docs.databricks.com/delta-live-tables>
  - **AutoLoader:** <https://docs.databricks.com/ingestion/auto-loader>
- 

**Congratulations!** You've completed a comprehensive, production-ready data engineering pipeline using Databricks Lakeflow Declarative Pipelines with automated incremental loading and full medallion architecture implementation.