



P4 - Raw

After running 1st and 2nd notebook you do the steps below:

1. login to databricks -> settings -> developer -> access token -> generate new -> copy it!



2. create key-vault



3. go to azvault-> new secret -> create it

Microsoft Azure

Home > Key vaults > Idbkv

Key vaults

Default Directory (nikhil0073outlook.onmicrosoft...)

+ Create

You are viewing a new version of Browse experience. Click here to access the old experience.

☐ Name 1

☐ dbkeyvault73

☒ Idbkv

Showing 1 - 2 of 2. Display auto

https://portal.azure.com/#

Idbkv | Secrets

Key vault

+ Generate/Import Refresh Restore Backup Manage deleted secrets View sample code

Name	Type	Status	Expiration date
RAGDB		✓ Enabled	

Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems
Access policies
Resource visualizer
Events
Objects
Keys
Secrets
Certificates
Settings
Monitoring
Automation
Help

Add or remove favorites by pressing Ctrl+Shift+F

Give feedback

Microsoft Azure

Home > Key vaults > Idbkv | Secrets

Create a secret

Upload options: Manual

Name *

Secret value *

Content type (optional):

Set activation date: ☒

Activation date: 11/26/2025 6:50:10 AM (UTC+05:30) Chennai, Kolkata, Mumbai, New Delhi

Set expiration date: ☐

Enabled: Yes No

Tags: 0 tags

Create Cancel

Remember the Name of the secret, and put the secret value that you copied in the databricks token



<https://adb-2321834658173720.0.azuredatabricks.net/#secrets/createScope>

Microsoft Azure databricks

modify the url as suggested...

Welcome to Databricks

Search data, notebooks, recent, and more... CTRL + P



3. Creating the Vector DB on Azure databricks

Databricks → Compute → Vector Search → Create your endpoint

enpoint_name = doc_vector_endpoint

index → select the table which contains the data which is docs_text

llm.rag.docs_idx → refer to the screenshots below to understand better.

4. Assign the Azure Key Vault user role to the **OID** (error occurs when we try to access the secret in notebook 3 -2nd cell). copy the **OID** from the ERROR.

Microsoft Azure databricks

Search data, notebooks, recents, and more... CTRL + P

ldb73-ws

Compute

All-purpose compute Job compute SQL warehouses **Vector Search** Pools Policies Apps Lakebase Postgres

Filter vector search endpoints by name Owned by me Empty endpoints Type

Status	Name	Type	Indexes	Creator
Provisioning	doc_vector_endpoint	Standard	0	admin@nikhil0073outlook.onmicrosoft.com

Just Provide the endpointname and craete it.

Microsoft Azure databricks

Search data, notebooks, recents, and more... CTRL + P

ldb73-ws

Catalog

Serverless Starter Wareho... Serverless S

Type to search...

My organization

- ldb73_ws
- system
- llm2
 - default
 - information_schema
 - rag
 - Tables(2)
 - docs_text**
 - docs_track
 - Volumes (1)
 - Models (1)
 - old-stuff
 - Delta Shares Received
 - samples
 - Legacy
 - hive_metastore

Catalog Explorer > llm2 > rag > **docs_text**

Open in a dashboard Share Create

Overview Sample Data Details Permissions Policies History Lineage Insights Quality

Description

AI generate Add

Filter columns...

Column	Type	Comment	Tags	Column maskir
id	bigint			
text	string			

About this table

Owner Databricks Admin

Type Managed

Open in a dashboard Share Create

Query

Notebook

Dashboard

Genie space

Vector search index

Synced table

Online table

Metric view

Create vector search index

Index structure

Name

llm2 rag docs_idx

Primary key

id

Columns to sync

text

Primary key and embedding column will automatically be synced to the index, even if they are not selected.

Embeddings

Embedding source

Compute embeddings Use existing embedding column

Embedding source column

text

Embedding model

databricks-gte-large-en

Cancel Create

Keep this exactly docs_idx

Create vector search index

Embedding source

Compute embeddings Use existing embedding column

Embedding source column

text

Embedding model

databricks-gte-large-en

Sync computed embeddings

Compute resources

Vector search endpoint

doc_vector_endpoint

Sync mode

Triggered Continuous

Advanced settings

Cancel Create

change the embedding model in the code also

adb-2321834658173720.0.azure.databricks.net/explore/data/llm2/rag/docs_idx?o=2321834658173720

select the idx and query it as below:

Query index

Request

```
{
  "num_results": 3,
  "columns": [
    "text"
  ],
  "query_text": "automatic features of the washer",
  "query_type": "HYBRID"
}
```

☒ Enable hybrid search
Combines vector and keyword search for better results.
[Learn more](#)

☐ Enable reranker
Reranks the top results with a secondary model for higher relevance.
[Learn more](#)

[Send request](#) [Reset example](#)

Response

```
{
  "manifest": {
    "column_count": 2,
    "columns": [
      {
        "name": "text"
      },
      {
        "name": "score"
      }
    ]
  },
  "result": {
    "row_count": 3,
    "data_array": [
      {
        "text": "temperature. clogged.\n While the washer fills, the water temperature may change as the automatic temperature control feature checks incoming water\ntemperature. This is normal.\n While the washer fills,"
      }
    ]
  }
}
```

About this vector index

State: Active

Owner: Databricks Admin [✎](#)

Type: Foreign

Data source: Vector Index

Popularity: [all](#)

Tags

[Add tags](#)

Insights

No active user in past 30 days

No tables joined in past 30 days

Related assets

[3 RAG Chatbot](#)

[All insights](#)

Microsoft Azure databricks

Search data, notebooks, recents, and more... CTRL + P

Workspace

2A Incremental PDF to docs_text **3 RAG Chatbot** 4 Chatbot GUI 1 Create needed tables

File Edit View Run Help Python Tabs: ON Last edit was now

2: Set needed parameters

```
host = "https://" + spark.conf.get("spark.databricks.workspaceUrl")
os.environ["DATABRICKS_TOKEN"] = dbutils.secrets.get(scope="P4-Scope", key="RAG08")
index_name = "llm2_rag.docs_idx"
host = "https://" + spark.conf.get("spark.databricks.workspaceUrl")
VECTOR_SEARCH_ENDPOINT_NAME = "doc_vector_endpoint"
```

3

38

4: Build Retriever

```
from databricks.vector_search.client import VectorSearchClient
from langchain_community.vectorstores import DatabricksVectorSearch
from langchain_community.embeddings import DatabricksEmbeddings

embedding_model = DatabricksEmbeddings(endpoint="databricks-gte-large-en")

def get_retriever(persist_dir: str = None):
    os.environ["DATABRICKS_HOST"] = host
    #Get the vector search index
```

Go to the Notebook 3 and run all the cells one-by-one.